


Claude GPT Qwen

EX1 (Serial CPU): per-benchmark best-of-N kernel speedup distribution



EX2 (OpenMP, 16 threads): per-benchmark best-of-N kernel speedup distribution

